

# SMT-based Compile-time Verification of Safety Properties for Smart Contracts

Leonardo Alt    Christian Reitwiessner

Ethereum Foundation

July 12, 2018

# Outline

- 1 Ethereum
- 2 Solidity BMC
- 3 Future Plans

# Outline

## 1 Ethereum

## 2 Solidity BMC

## 3 Future Plans

# Ethereum

- Global networked application platform
- Distributed public database
- Blockchain consensus

# Ethereum

- Global networked application platform
- Distributed public database
- Blockchain consensus
- Trustless
- Transparent
- No single entity has control

# Smart Contracts

- Accounts controlled by code
- SCs have a storage which can only be written by the SC itself
- The code runs on the Ethereum Virtual Machine (EVM)
- Opcodes have costs (*gas*)
- Transactions may *revert* (exception termination state)

# Smart Contracts

- The code is public and immutable after creation
- Monetary incentive to study/attack the program
- The DAO and Parity hacks/accidents

# Smart Contracts

- The code is public and immutable after creation
- Monetary incentive to study/attack the program
- The DAO and Parity hacks/accidents

Easy, just use formal verification :)



# Smart Contracts - Verification Frameworks

- EVM Formal Semantics: Eth-Isabelle [8], KEVM [9], Ethereum-Lem [10]
- EVM bytecode Symbolic Execution: Oyente [5], Mythril [3], MAIAN [6]
- Translation of Solidity to verifiable languages: Why3 [7], F\* [2], ZEUS [4]
- Our approach: Solidity SMT-based BMC
  - Part of the compilation stack
  - No extra effort spent by the developer
  - Counterexamples

# Outline

1 Ethereum

2 Solidity BMC

3 Future Plans

# Solidity

- Developed for smart contracts
- Most used language to write smart contracts
- Syntax similar to C/Java
- Main code elements are *contracts*
- Storage and local variables
- Integers of various sizes, *address*, *mapping*, structs, arrays...

# Solidity

```
1  contract Token {
2      /// The main balances / accounting mapping.
3      mapping(address => uint256) balances;
4
5      /// Create the token contract crediting 'msg.sender' with
6      /// 10000 tokens.
7      constructor() public {
8          balances[msg.sender] = 10000;
9      }
10
11     /// Transfer '_value' tokens from 'msg.sender' to '_to'.
12     function transfer(address _to, uint256 _value) public {
13         require(balances[msg.sender] >= _value);
14         uint256 sumBefore = balances[msg.sender] + balances[_to];
15         balances[msg.sender] -= _value;
16         balances[_to] += _value;
17         uint256 sumAfter = balances[msg.sender] + balances[_to];
18         assert(sumBefore == sumAfter);
19     }
20 }
```

# Solidity BMC

- Basic idea: SMT-based BMC
- Gas incentive against unbounded loops
- Practicality over Completeness
- Precise encoding

# Solidity BMC

Verification targets:

- Underflow / Overflow / Division by 0
- Trivial conditions / Unreachable code
- Assertions

# Solidity → SMT Encoding

- Traverse the AST
  - Collect constraints
  - Query the SMT solver
- Five types of constraints:  
Control-flow, Type constraint, Variable assignment, Branch conditions, Verification Target

# Solidity → SMT Encoding

Branch conditions:

Auxiliary stack that keeps track of the conjunction of conditions that are true at the current program path.

No constraint is added to the solver.



# Solidity $\rightarrow$ SMT Encoding

Control-flow:

Global constraints of the form  $b \rightarrow r$ , where  $b$  is the conjunction of conditions in the current path and  $r$  is the condition in `require(r)` and `assert(r)`.

# Solidity → SMT Encoding

Type constraint:

Declaration of local variables use default values from types (Integer: 0, Bool: false), and function parameters are initialized with a range of valid values (`uint32 x`:  $0 \leq x < 2^{32}$ ).

# Solidity → SMT Encoding

Variable assignment:

The encoding follows the SSA form. When a variable is assigned inside different branches, a new `if-then-else` variable is created after the branching to re-combine the different values.

# Solidity → SMT Encoding

Verification target:

- Underflow / Overflow / Division by 0
- Trivial conditions / Unreachable code
- Assertions

Counterexamples are given when a property fails.

# Solidity $\rightarrow$ SMT Encoding

```
1  contract C
2  {
3      function f(uint256 a, uint256 b)
4      {
5          if (a == 0)
6              require(b <= 100);
7          else if (a == 1)
8              b = 1000;
9          else
10             b = 10000;
11         assert(b <= 100000);
12     }
13 }
```

1.  $a_0 \geq 0 \wedge a_0 < 2^{256} \wedge$
2.  $b_0 \geq 0 \wedge b_0 < 2^{256} \wedge$
3.  $(a_0 = 0) \rightarrow (b_0 \leq 100) \wedge$
4.  $b_1 = 1000 \wedge b_2 = 10000$
5.  $b_3 = \text{ite}(a == 1, b_1, b_2) \wedge$
6.  $b_4 = \text{ite}(a == 0, b_0, b_3) \wedge$
7.  $\neg b_4 \leq 100000$

# Outline

1 Ethereum

2 Solidity BMC

3 Future Plans

# Future Plans - Loops

- Automatic detection of loop bounds
- Invariant annotations
- Current syntax:

```
1  contract C
2  {
3      uint x;
4      function f(uint[] _arr) {
5          require(_arr.length < 12);
6          for (uint i = 0; i < _arr.length; ++i)
7              x += _arr[i];
8      }
9  }
```

# Future Plans - Multi-transaction Invariants

```
1  contract C
2  {
3      uint a;
4
5      constructor () public {}
6
7      function a1() public { a = 1; }
8      function a2() public { a = 2; }
9      function a3() public { a = 3; }
10     function a4() public { a = 4; }
11
12     function plusA(uint x) public view returns (uint) {
13         require(x < 1000);
14         return a + x;
15     }
16 }
```

No transaction leads to  $a > 4$



# Future Plans - Post-constructor Invariants

```
1  contract C
2  {
3      uint a;
4
5      constructor (uint x) public {
6          require(x < 10);
7          a = x;
8      }
9
10     function plusA(uint x) public view returns (uint) {
11         require(x < 1000);
12         return a + x;
13     }
14 }
```

State variable *a* can only be assigned values smaller than 10, and is never assigned again.

# Future Plans - Modifiers as pre and postconditions

```
1  contract C
2  {
3      uint totalSupply;
4
5      modifier safeBalance {
6          require(totalSupply == 10000);
7          _;
8          assert(totalSupply == 10000);
9      }
10
11     function transfer(address _to, uint256 _value) safeBalance {
12         ...
13     }
14 }
```

Modifier `safeBalance` guarantees that the `totalSupply` does not change after a transfer.

# Future Plans - Function abstraction

```
1  contract C
2  {
3      uint totalSupply;
4
5      modifier safeBalance {
6          require(totalSupply == 10000);
7          -;
8          assert(totalSupply == 10000);
9      }
10
11     function transfer(address _to, uint256 _value) safeBalance {
12         ...
13     }
14
15     function zeroAccount(address _to) {
16         transfer(_to, balance[msg.sender]);
17         assert(totalSupply == 10000);
18     }
19 }
```

The modifier can then be used to abstract function calls and prove properties more efficiently.

# Future Plans - Range restriction of real life values

- `block.timestamp` will not exceed 64 bits for the next 500 billion years
- `block.number`

# Future Plans - Effective Callback Freeness

- Recently introduced by [1]
- State variable invariants still hold after external calls

# Let's build together

`github.com/ethereum/solidity`

`gitter.im/ethereum/solidity-dev`

`leo@ethereum.org`

Thank you!

# References I

- [1] Grossman, S., Abraham, I., Golan-Gueta, G., Michalevsky, Y., Rinetzky, N., Sagiv, M., Zohar, Y.: Online detection of effectively callback free objects with applications to smart contracts. Proc. ACM Program. Lang. 2(POPL), 48:1–48:28 (2017)
- [2] Bhargavan, K., Delignat-Lavaud, A., Fournet, C., Gollamudi, A., Gonthier, G., Kobeissi, N., Kulatova, N., Rastogi, A., Sibut-Pinote, T., Swamy, N., Zanella-Béguelin, S.: Formal verification of smart contracts: Short paper. In: Proceedings of the 2016 ACM Workshop on Programming Languages and Analysis for Security. pp. 91–96. PLAS '16 (2016)

# References II

- [3] ConsenSys: Mythril (2018),  
[github.com/ConsenSys/mythril](https://github.com/ConsenSys/mythril)
- [4] Kalra, S., Goel, S., Dhawan, M., Sharma, S.: ZEUS: Analyzing safety of smart contracts (2018)
- [5] Luu, L., Chu, D.H., Olickel, H., Saxena, P., Hobor, A.: Making smart contracts smarter. In: Proceedings of the 2016 ACM SIGSAC Conference on Computer and Communications Security. pp. 254–269. CCS '16 (2016)



# References III

- [6] Nikolic, I., Kolluri, A., Sergey, I., Saxena, P., Hobor, A.:  
Finding the greedy, prodigal, and suicidal contracts at scale.  
CoRR abs/1802.06038 (2018),  
<http://arxiv.org/abs/1802.06038>
- [7] Why3: Why3 (2018), [why3.lri.fr](http://why3.lri.fr)
- [8] Eth-Isabelle, [github.com/pirapira/eth-isabelle](https://github.com/pirapira/eth-isabelle)
- [9] KEVM, [github.com/kframework/evm-semantics](https://github.com/kframework/evm-semantics)
- [10] Ethereum-Lem, [github.com/mrsmkl/ethereum-lem](https://github.com/mrsmkl/ethereum-lem)